

Latency (ms, log scale)

10^5
 10^4
 10^3
 10^2

Reflexer median is 11.5× faster than Thinker

p95 1130ms

med 395ms

Reflexer
(SLM, n=35)

med 4556ms

Thinker
(LLM, n=38)

